

Machine Learning Application:

Diagnosing Pneumonia With Chest X-Ray Images

Eshaan Vora, Jeremy Anderson, Nima Nakhjavani

Note:

We ran our neural net in a GPU environment using NVIDIA CUDA parallel computing toolkit and Tensorflow GPU. We used data augmentation during training, extending the conv_base (vgg16) model and running it end to end on the inputs which is an extremely expensive technique, intractable on the CPU.

Final Model Results:

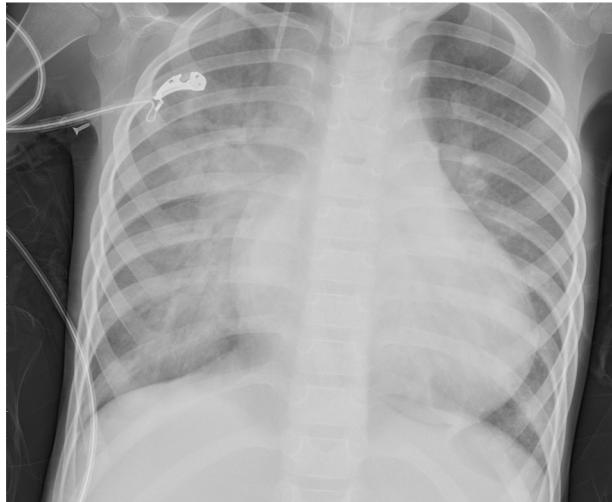
test accuracy: 0.971875

test loss: 0.0031179

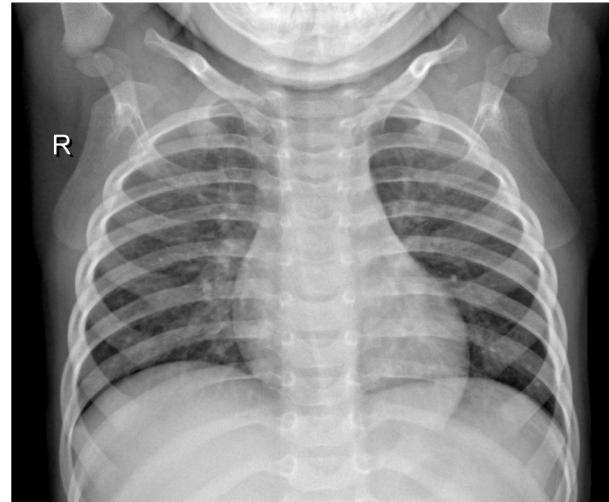
Objective:

Determine if a hospital patient has pneumonia based off of their chest x-ray images

Pneumonia



Healthy



Our Dataset

Organized into three folders

test , train, and validation

Sub-folders: pneumonia and normal

Consists of over 5,000 chest x-ray images

- Both healthy and pneumonia positive lungs

Chest x-rays were selected from pediatric patients 1-5 years old

All x-rays were performed as part of routine care

We created our own test train and validation subsets within the dataset

Allowed us to train our models more accurately

Why this Relevant?

- Pneumonia is present in many severe COVID-19 cases - can be deadly in vulnerable populations such as the elderly
- Real world application
 - Could help take strain off medical professionals, providing an initial diagnostic evaluation focused on minimizing false negatives
 - Assist in more accurately predicting pneumonia cases
 - Can process and predict a large number of x-rays at a much more efficient rate
- Could be used to determine other lung conditions / other diseases in general
- Since x-rays were done as part of routine care, models can be used to spot these issues before patients even know

Our Final Model

- VGG16 - pre trained convolutional neural network model for large scale image recognition (Transfer Learning)
 - Added a convolutional base and then flattened to an output with a dimensionality of 256 - used relu activation
 - Ran dropout with a rate of 0.35
 - Flattened to final output with dimensionality of 1 - used sigmoid function because it is a binary classification (pneumonia or health)
 - Close to 17 million trainable parameters by the end
- Data augmentation allowed us to increase the amount of data available for the training model without having to physically add new data by scaling and rotating pre-existing images
- Fine tuned our model by freezing the convolutional base to get our accuracy up even higher

Model Images

```
Epoch 86/100
33/33 [=====] - 44s 1s/step - loss: 0.0478 - acc: 0.9788 - val_loss: 0.0222 -
val_acc: 0.9531
Epoch 87/100
33/33 [=====] - 46s 1s/step - loss: 0.0768 - acc: 0.9712 - val_loss: 1.0973 -
val_acc: 0.9125
Epoch 88/100
33/33 [=====] - 45s 1s/step - loss: 0.0545 - acc: 0.9773 - val_loss: 0.1436 -
val_acc: 0.9258
Epoch 89/100
33/33 [=====] - 43s 1s/step - loss: 0.0795 - acc: 0.9712 - val_loss: 0.1811 -
val_acc: 0.9375
Epoch 90/100
33/33 [=====] - 43s 1s/step - loss: 0.0674 - acc: 0.9697 - val_loss: 0.1999 -
val_acc: 0.9194
Epoch 91/100
33/33 [=====] - 43s 1s/step - loss: 0.0717 - acc: 0.9742 - val_loss: 0.1186 -
val_acc: 0.9312
Epoch 92/100
33/33 [=====] - 43s 1s/step - loss: 0.0467 - acc: 0.9833 - val_loss: 0.3639 -
val_acc: 0.9355
Epoch 93/100
33/33 [=====] - 44s 1s/step - loss: 0.0662 - acc: 0.9712 - val_loss: 0.0296 -
val_acc: 0.9344
Epoch 94/100
33/33 [=====] - 44s 1s/step - loss: 0.0503 - acc: 0.9833 - val_loss: 1.0079 -
val_acc: 0.9000
Epoch 95/100
33/33 [=====] - 44s 1s/step - loss: 0.0568 - acc: 0.9773 - val_loss: 0.5223 -
val_acc: 0.9094
Epoch 96/100
33/33 [=====] - 45s 1s/step - loss: 0.0847 - acc: 0.9652 - val_loss: 0.0737 -
val_acc: 0.9290
Epoch 97/100
33/33 [=====] - 44s 1s/step - loss: 0.0610 - acc: 0.9833 - val_loss: 0.0732 -
val_acc: 0.9250
Epoch 98/100
33/33 [=====] - 45s 1s/step - loss: 0.0656 - acc: 0.9742 - val_loss: 0.4918 -
val_acc: 0.9419
Epoch 99/100
33/33 [=====] - 45s 1s/step - loss: 0.0492 - acc: 0.9803 - val_loss: 0.2409 -
val_acc: 0.8875
Epoch 100/100
33/33 [=====] - 44s 1s/step - loss: 0.0590 - acc: 0.9727 - val_loss: 0.2595 -
val_acc: 0.9226
```

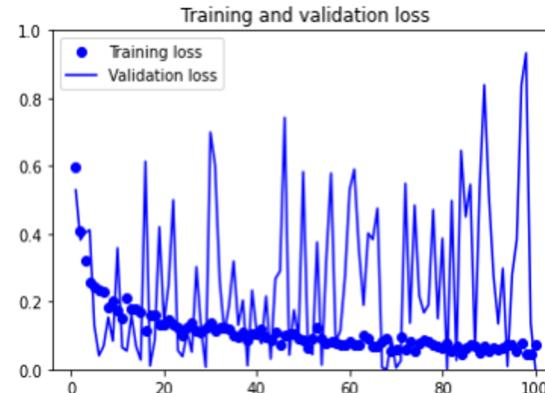
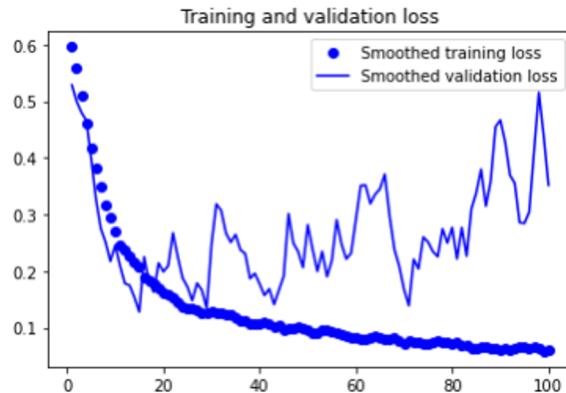
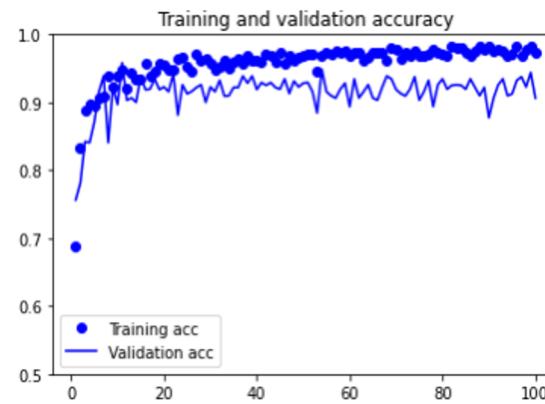
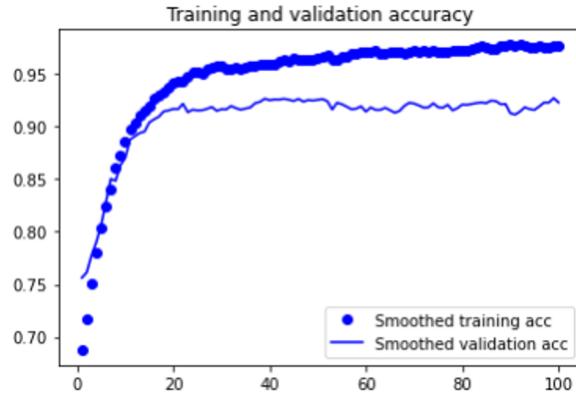
Convolutional
Base:

Model: "vgg16"		
Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0

Model
Summary:

Model Summary		
Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_5 (Flatten)	(None, 8192)	0
dense_9 (Dense)	(None, 256)	2097408
dropout_2 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 1)	257
Total params: 16,812,353		
Trainable params: 16,812,353		
Non-trainable params: 0		

Model Results



test acc: 0.971875011920929
test loss: 0.0031179096549749374

Future Steps

- Reevaluate neural network architecture and readjust to improve accuracy, loss, and/or runtime
 - Even on a GPU, this algorithm is very expensive, so goal is to experiment with the architecture improve runtime without compromising model performance
 - Ex: Trying various dropout rates and learning rates
- Determining the precision and recall of the algorithm
- Implement K-Fold Cross Validation
- Testing model with different image datasets

Dataset Link

<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>